

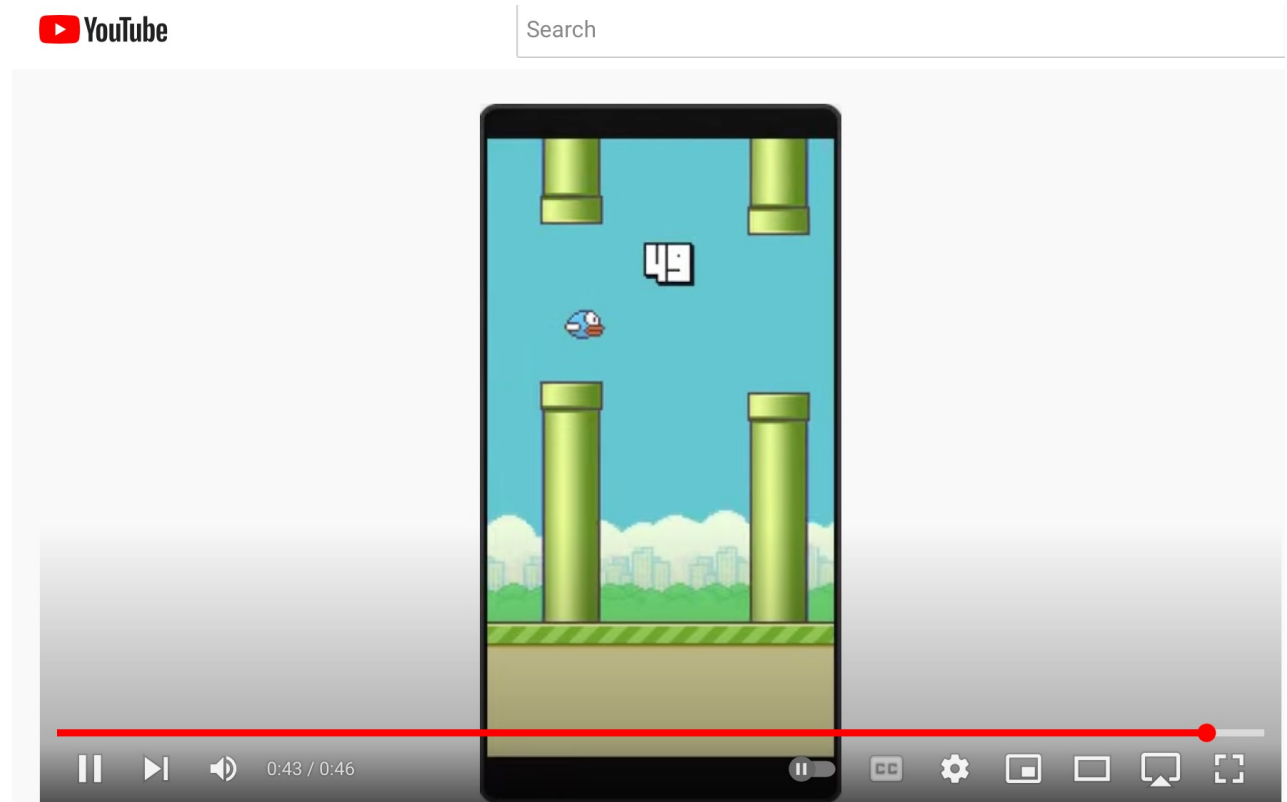
# Overview of Reinforcement Learning

Dr. Lotfi ben Othmane  
University of North Texas

# The Flappy Bird Game

<https://www.youtube.com/watch?v=BV7a4rufMOg&t=5s>

Can you do it with  
supervised  
machine learning?

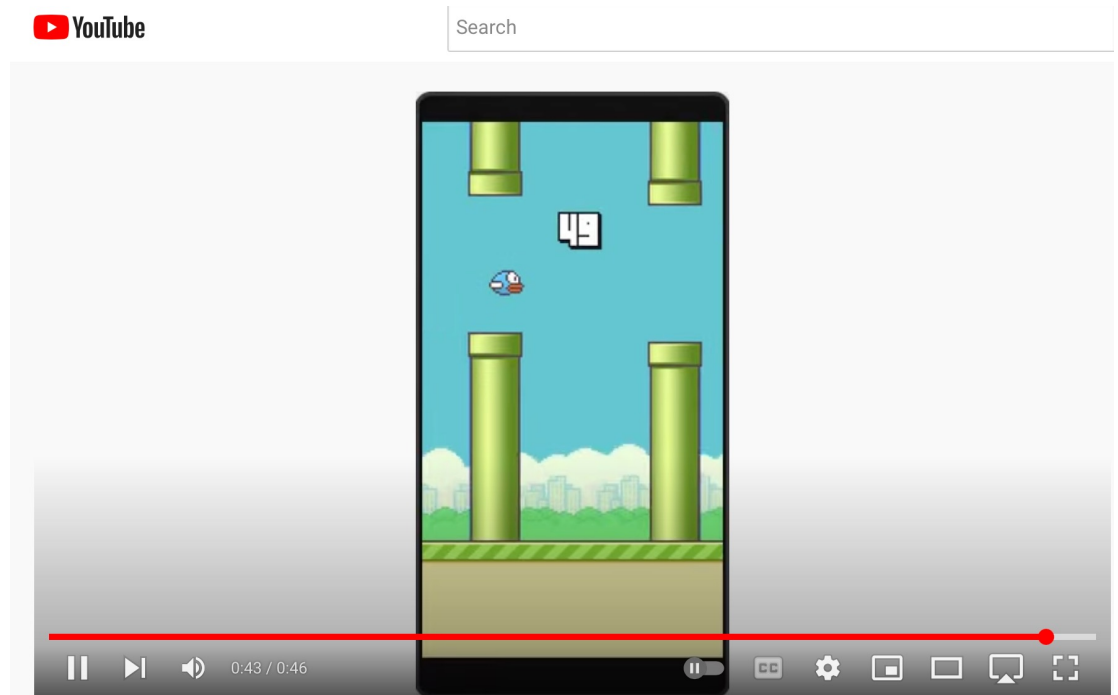


# The Flappy Bird Game

## Reward

**+1** if Flappy Bird is still **alive**

**-1000** if Flappy Bird is **dead**



# The Flappy Bird Game Program

<http://sarvagyaish.github.io/FlappyBirdRL/>

## Inputs

$S$  is a set of states  
 $A$  is a set of actions  
 $\gamma$  the discount  
 $\alpha$  is the step size

## Local

real array  $Q[S,A]$   
previous state  $s$   
previous action  $a$   
initialize  $Q[S,A]$  arbitrarily  
observe current state  $s$   
**repeat**

    select and carry out an action  $a$   
    observe reward  $r$  and state  $s'$   
     $Q[s,a] \leftarrow Q[s,a] + \alpha(r + \gamma \max_{a'} Q[s',a'] - Q[s,a])$   
     $s \leftarrow s'$  **until** termination

**S**: State as the step

**A**: Actions are: (1) jump or  
(2) do nothing

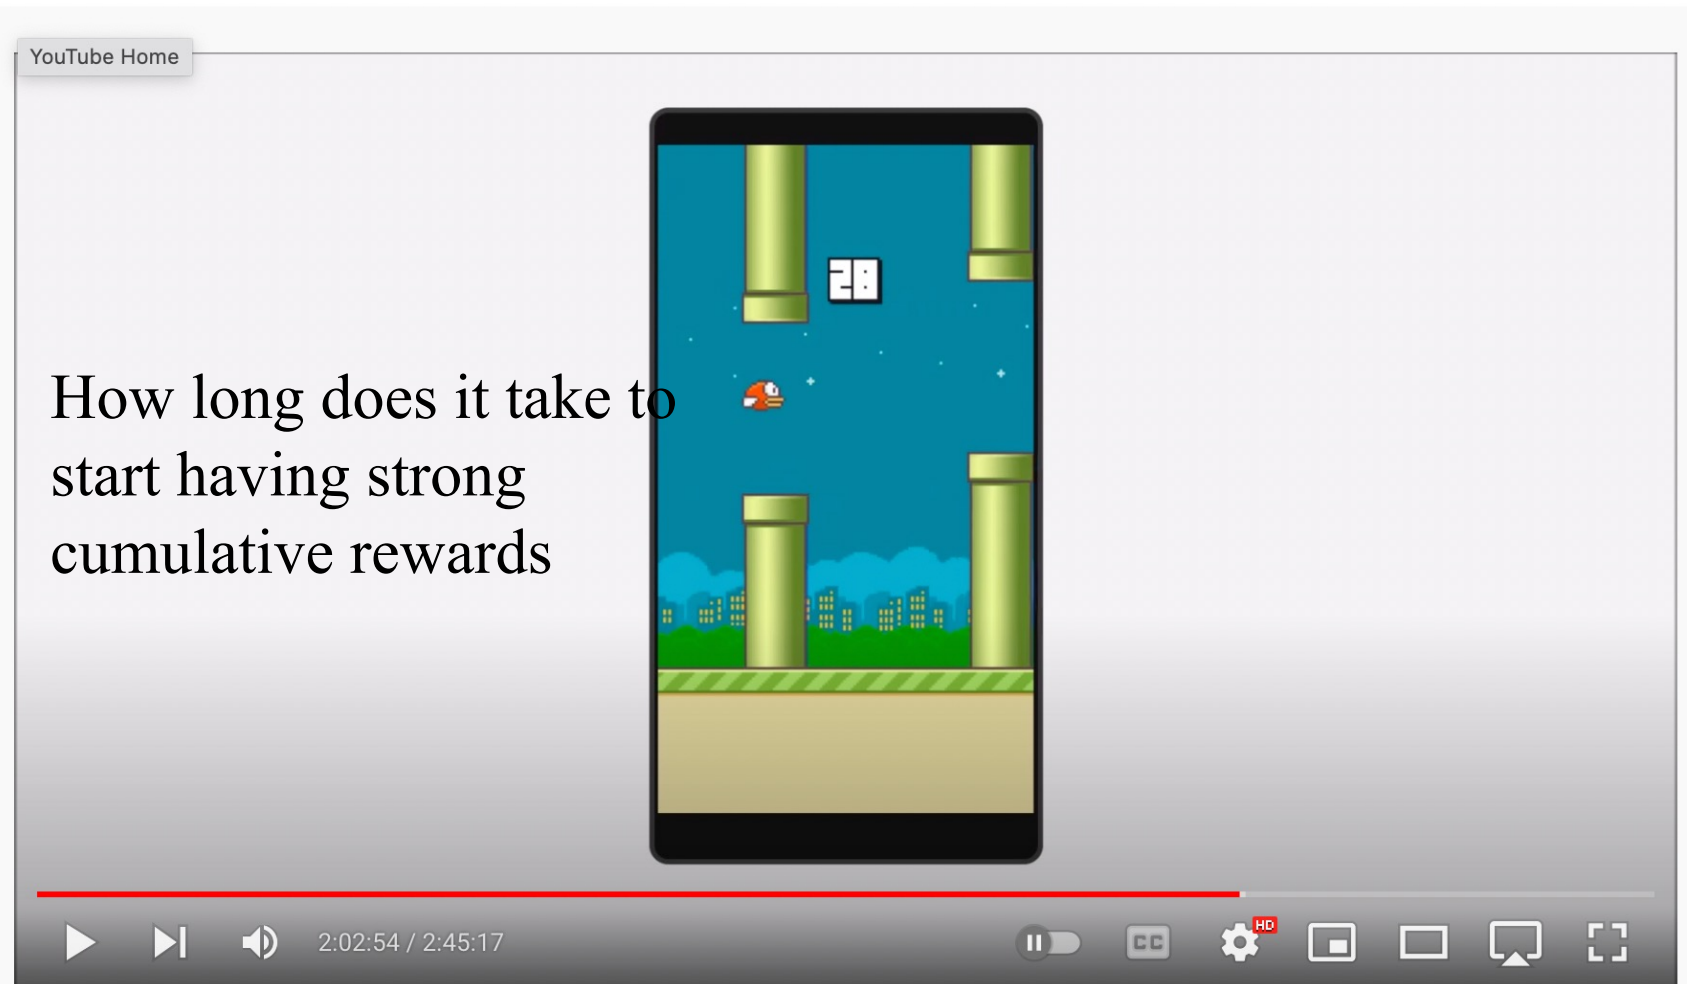
**$\alpha$** : Learning rate— value 0.7

**r**: Reward +1 or -1000

**$\gamma$** : Discount factor (set to 1)

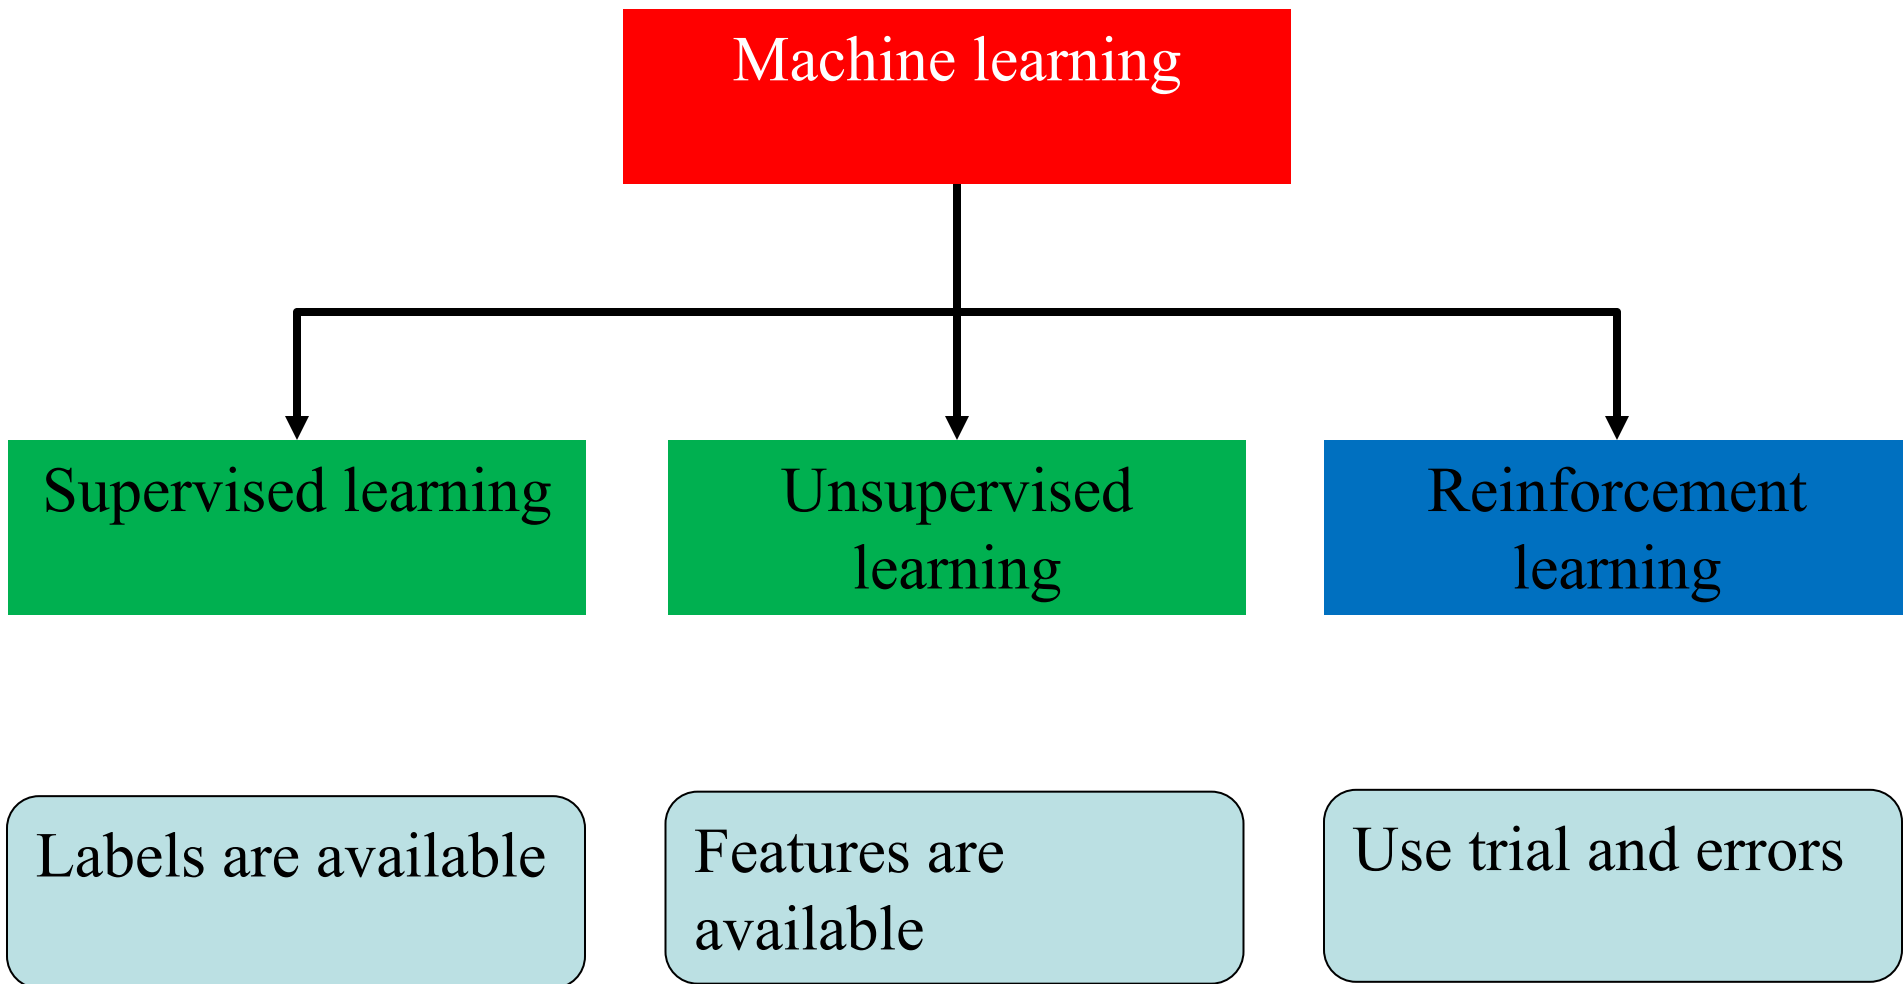
What about the function Q learning?

# The Flappy Bird Game – Q Learning Over Time



<https://www.youtube.com/watch?v=OJw4HTWvGdY&t=21s>

# Reinforcement Learning




# Reinforcement Learning

- In a lot of cases, like in games, the data is not good enough to have an intelligent machine. But, we have knowledge on how to make decisions.

# Reinforcement Learning

**Reinforcement learning** is the problem faced by an agent that must learn behavior through trial-and-error interactions with dynamic environment.



<https://en.wikipedia.org/wiki/Reinforcement>



# Reinforcement Learning

The agent **interacts** with its **environment** in **discrete time steps**. At each time  **$t$** , the agent receives the current state  $s_t$  and reward  $r_t$ . It chooses an action  $a_t$  from the set of available actions, which is subsequently sent to the **environment**. The environment moves to a new state  $s_{t+1}$  and the reward  $r_{t+1}$  associated with the **transition**  $(s_t, a_t, s_{t+1})$  is determined.

In a policy RL, the goal of a reinforcement learning agent is to learn a **policy**

$\pi: A \times S \rightarrow [0, 1]$ ,  $\pi(a, s) = \Pr(a_t = a | s_t = s)$  which **maximizes the expected cumulative reward**.

# Reinforcement Learning

Reinforcement learning could be modeled as a **Markov Decision Process**.

**S**: Set of environment and agent states

**A**: Set of actions of the agent

$P_a(s, s') = P(s_{t+1} = s' | s_t = s, a_t = a)$  is the probability of transition (at time **t**) from state **s** to state **s'** under action **a**.

$R_a(s, s')$  is the immediate reward after transition from **s** to **s'** with action **a**.

Goal: Maximize the cumulative reward  $R_a(s, s')$

# Simple Approach

Brute force approach

1. For each possible policy, sample returns while following it
2. Choose the policy with the largest expected return

One problem with this is that the number of policies can be large, or even infinite.

# RL Exploration vs. Exploitation

- Good decisions are through exploitation of knowledge  
=> Pick the action that has the highest utility
- Exploration: Pick a randomly reasonably good action
- The utility weights the sub-optimal exploitation decision with  $0 < \epsilon < 1$

# Applications of RL

- There many applications of RL including:
  1. Games
  2. Robotics and control

# Exercise 1

- Formulate cleaning robot using reinforcement learning
  - What is the rewards?
  - What are the states?
  - What are the transitions?

# Exercise 2

- Formulate the robot using reinforcement learning
  - What is the rewards?
  - What are the states?
  - What are the transitions?

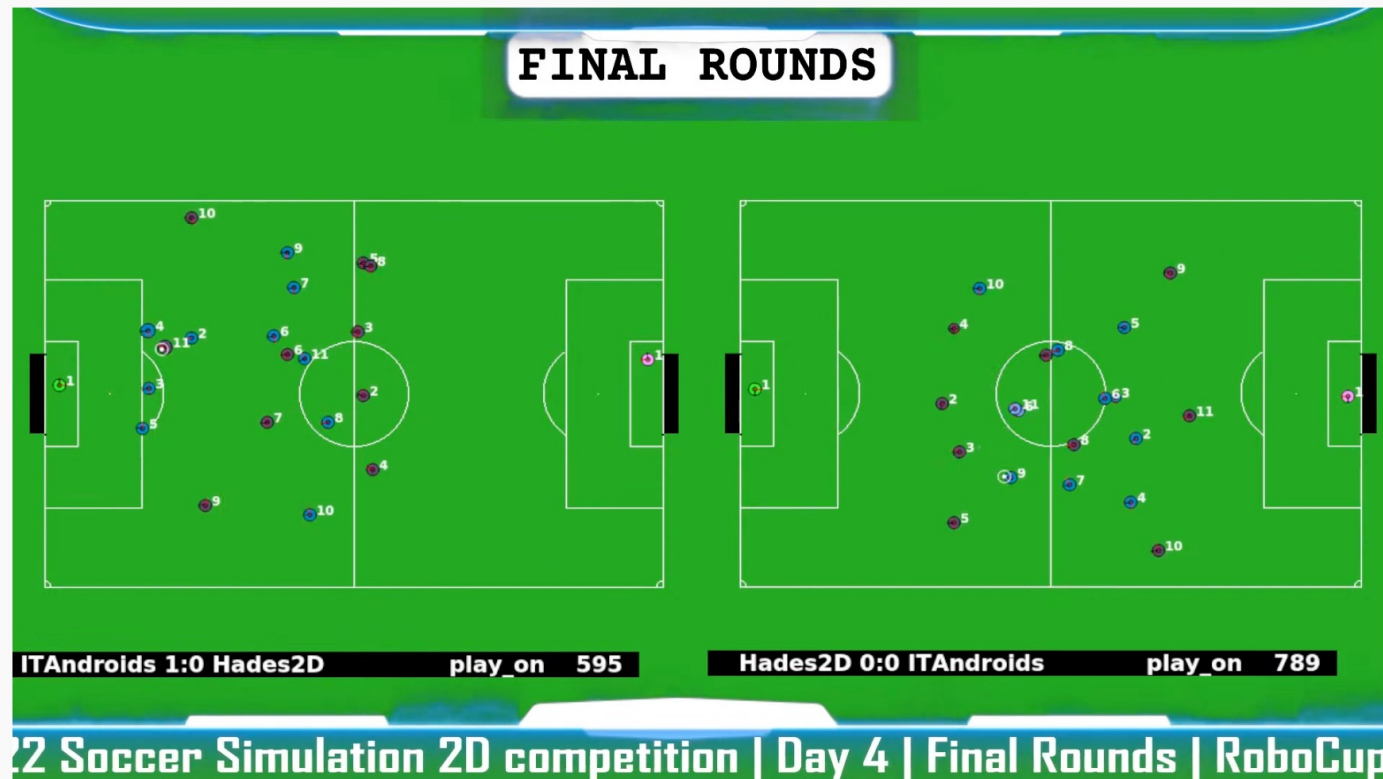


# RL for RoboCup Soccer

[https://www.youtube.com/watch?v=51L\\_gq\\_fu5U](https://www.youtube.com/watch?v=51L_gq_fu5U)



Search





# RL with Unknown Reward

- Receive a reward **when** the agent puts the ball in the net.
- The agent needs to catch the ball, give the ball to team members, advance with the ball, and shoot the ball.
- How RL can help in such scenarios

<https://rcsoccersim.github.io>

<https://github.com/rcsoccersim/>

## The RoboCup Soccer Simulator


### About the RoboCup Soccer Simulator

The RoboCup Soccer Simulator is a research and educational tool for multiagent systems and artificial intelligence. It enables for two teams of 11 simulated autonomous robotic players to play soccer (football).

### League Overview

Without the necessity to maintain any robot hardware, the RoboCup Simulation League's focus comprises artificial intelligence and team strategy.

# Learning Agents



# Summary

- **Reinforcement learning** is the problem faced by an agent that must learn behavior through trial-and-error interactions with dynamic environment.
- Reinforcement learning methods looks to optimize the reward from the decisions.
- It has many applications mostly in games and robotics/control theory

Thank you

Any Question?